

Introduction to Machine Learning with Python

Regularization

Instructor: Tugce Gurbuz

Aug 10th, 2023

Complex Models Need Regularization

In ML, we use highly complex models with many adjustable parameters.

- Learning the noise in data?




Model Complexity Matters!

- Too simple -> “underfits”



Model Complexity Matters!


- Too complex -> “overfits”



Model Complexity Matters!

Bias vs Variance Tradeoff



- Bias: The simple models often give systematically too small weights.



Model Complexity Matters!


Bias vs Variance Tradeoff

- Bias: The simple models often give systematically too small weights.
- Variance: The complex model capture the variance too much that leads to poor generalization



Model Complexity Matters!

How to pick right model complexity?



Model Complexity Matters!

The magic!

Deep learning on images (Zhang, Bengio, Hardt, Recht, and Vinyalsn 2017)

- gives 0 training error -- and small test error
- gives 0 training error with randomized labels

Model Complexity Matters!

The magic!

Deep learning on images (Zhang, Bengio, Hardt, Recht, and Vinyalsn 2017)

- gives 0 training error -- and small test error
- gives 0 training error with randomized labels

Large language models memorize a lot

- GPT-3 (175B params trained on 500B words) seems to memorize a lot

Q. What do you call a droid that takes the long way around?

Model Complexity Matters!

The magic!

Deep learning on images (Zhang, Bengio, Hardt, Recht, and Vinyalsn 2017)

- gives 0 training error -- and small test error
- gives 0 training error with randomized labels

Large language models memorize a lot


- GPT-3 (175B params trained on 500B words) seems to memorize a lot

Q. What do you call a droid that takes the long way around?
R2 detour.

Characteristics of Regularized Models

The more regularized models give us:

- Smaller weights (less fitting to noise)
- Smoother models
- Models with lower capacity



Common Regularization Techniques

L1 and L2 penalties:

- Train to minimize normal loss + $c * L1(\text{weights})$
 - L1: Lasso regression
 - Drives some weights to 0

Common Regularization Techniques



L1 and L2 penalties:

- Train to minimize normal loss + $c * L1(\text{weights})$
 - L1: Lasso regression
 - Drives some weights to 0
- Train to minimize normal loss + $c * L2(\text{weights})$
 - L2: ridge regression
 - Makes biggest weights smaller.

Common Regularization Techniques

L1 and L2 penalties:

- Train to minimize normal loss + $c * L1(\text{weights})$
 - L1: Lasso regression
 - Drives some weights to 0
- Train to minimize normal loss + $c * L2(\text{weights})$
 - L2: ridge regression
 - Makes biggest weights smaller.



Common Regularization Techniques

L1 and L2 penalties:


- Train to minimize normal loss + $c * L1(\text{weights})$
 - L1: Lasso regression
 - Drives some weights to 0
- Train to minimize normal loss + $c * L2(\text{weights})$
 - L2: ridge regression
 - Makes biggest weights smaller.

L-infinity penalty:


- Train to minimize normal loss - but don't let the weights get too big

Common Regularization Techniques

Drop-out



(a) Standard Neural Net



(b) After applying dropout.

Common Regularization Techniques

Early stopping -> Avoid overfitting by stopping your training at the right time

Common Regularization Techniques

Data Augmentation




Image Augmentation Sample. Image by <https://github.com/aleju/imgaug>

Common Regularization Techniques


Stochastic Gradient Descent (SGD)

- Initialize with small random weights
- Weights get bigger as one iterates
- Use early stopping to avoid overfitting


Common Regularization Techniques

Stochastic Gradient Descent (SGD)

- Initialize with small random weights
- Weights get bigger as one iterates
- Use early stopping to avoid overfitting



(a) CIFAR-10 (Train)



(b) CIFAR-10 (Test)

Common Regularization Techniques

Learning rates and regularization

- Small learning rates are more likely to find a deep minimum -> might be good or bad

Common Regularization Techniques

Learning rates and regularization

- Small learning rates are more likely to find a deep minimum -> might be good or bad
- Large learning rates misses deep minimas, finds broader and flatter minimas that may be more robust.

Common Regularization Techniques

Conclusion on regularization techniques: Use all!

Common Regularization Techniques

Hyperparameter tuning is a search:

- **Grid Search:** Try all possible combinations of hyperparameters
- **Random Search:** Randomly try different combinations of hyperparameters
- **Coordinate-wise Gradient Descent:** Start at one set of hyperparameters and try changing one at a time, accept any changes that reduce your validation error
- **Bayesian Optimization / Auto ML:** Start from a set of hyperparameters that have worked well on a similar problem, and then do some sort of local exploration (e.g., gradient descent) from there.